

Attention-Driven Dropout

A Simple Method to Improve Self-supervised Contrastive Sentence Embeddings

**Fabian Stermann^{1,2}, Ilias Chalkidis⁴, Amihossein Vahidi^{1,3},
Bernd Bischl^{1,3}, Mina Rezaei^{1,3}**

1. Department of Statistics, LMU Munich, Munich, Germany

2. Convalid Analytics, Munich, Germany

3. Munich Center for Machine Learning, Munich, Germany

4. Department of Computer Science, University of Copenhagen, Denmark



- Contrastive Learning

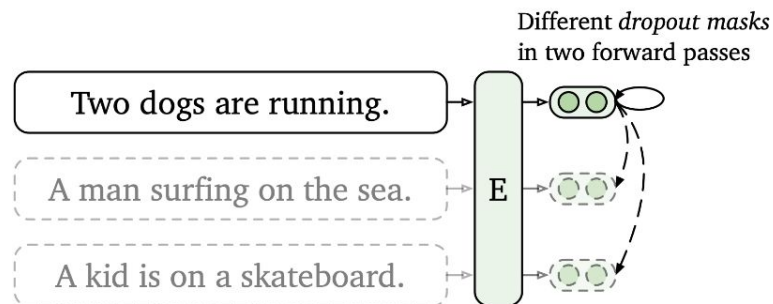
- **Positive** examples: similar sentences
- **Negative** examples: dissimilar sentences
- Objective
 - Minimize distance of positive examples
 - Maximize distance of negative examples

- **Unsupervised SimCSE** (Gao et al. 2021)

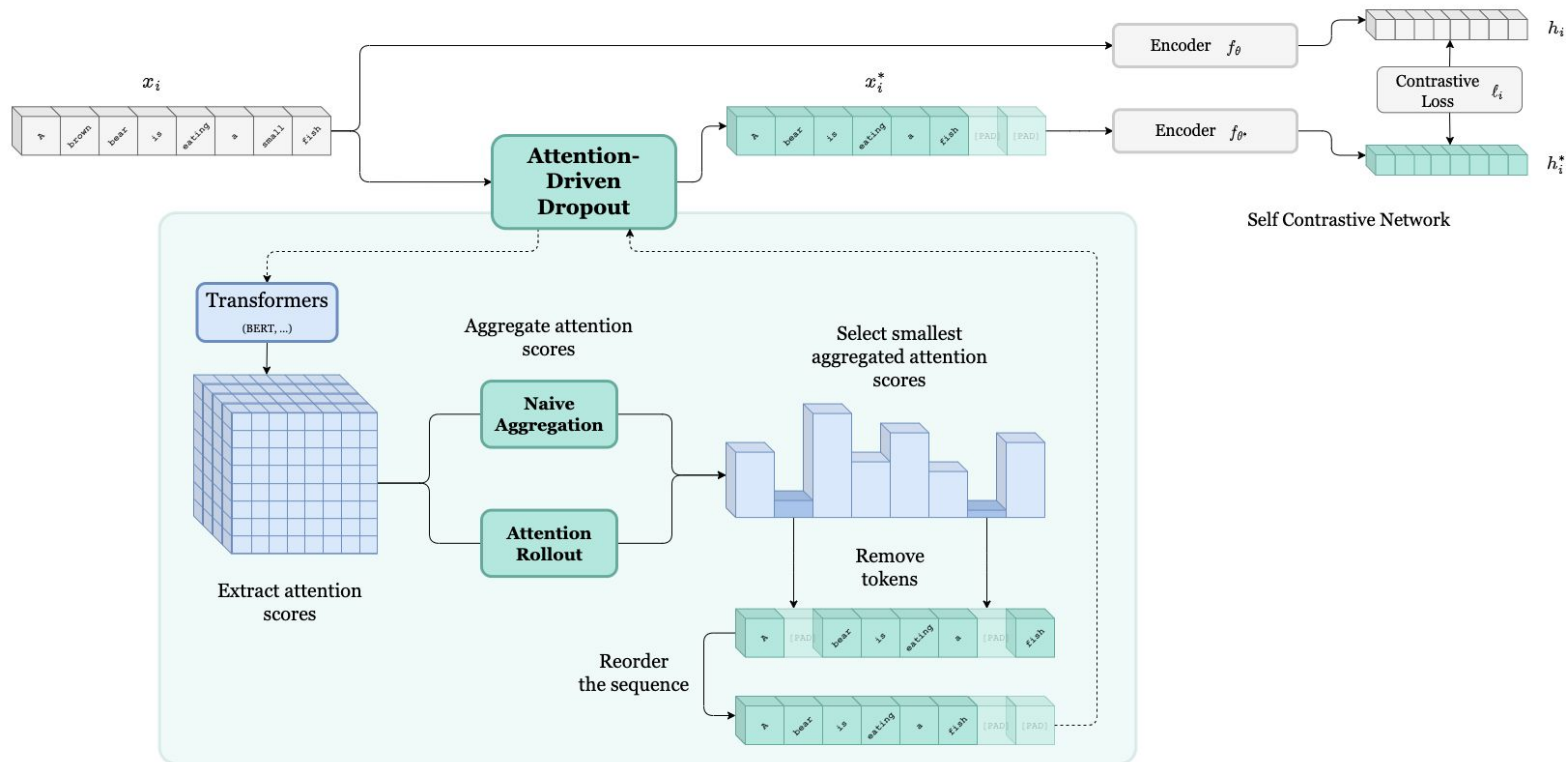
- Positive instances simply have different dropout masks
- Uses in-batch samples as negatives

$$\ell_{(x_i, x_j)} = -\log \frac{\exp(\text{sim}(\mathbf{h}_i, \mathbf{h}_j)/\tau)}{\sum_{k=1}^{2N} \mathbb{I}_{[k \neq i]} \exp(\text{sim}(\mathbf{h}_i, \mathbf{h}_k)/\tau)}$$

(a) Unsupervised SimCSE



- E** Encoder
→ Positive instance
- → Negative instance



- Naive Aggregation

- Summation across layers and heads
- “Attentiveness” of tokens to specific tokens

$$a_i = \sum_{l=1}^L \sum_{h=1}^H \sum_{s=1}^{S_1} A_{lihs}$$

- Rollout Aggregation

- *Abnar and Zuidema (2020)*
- Approximate Attention to input tokens

$$\tilde{A}(l) = \begin{cases} A(l)\tilde{A}(l^-) & \text{if } l > l^- \\ A(l) & \text{if } l = l^- \end{cases}$$

$$a_i = \sum_{l=1}^L \sum_{s=1}^{S_1} \tilde{A}_{lis}$$

Naive Aggregation Example


0: A brown bear is eating a small fish.

1: A brown bear is eating a ~~small~~ fish.

2: A ~~brown~~ bear is eating a ~~small~~ fish.

3: A ~~brown~~ bear is eating ~~a~~ ~~small~~ fish.

4: A ~~brown~~ bear is eating ~~a~~ ~~small~~ ~~fish~~.



Static

- Use a predefined constant k as the dropout rate.
- Fixed token removal across sequences, regardless of length.


Dynamic

- Remove more redundant information in long sequence.
- k is calculated dynamically based on non-padding tokens.
- For sequence x_i^+ with t_s tokens:


$$k = \left\lfloor \frac{t_s}{t} \right\rfloor$$

- Select $g_i = \text{imin}(a_i, k) \in \mathbb{R}^k$, the indices of k lowest attention scores.
- Replace x_{ij}^+ with padding for $j \in g_i$, forming x_i^* .
- Reorder x_i^* with padding tokens aligned to the right.

Naive



Rollout



k Number of tokens to drop

Dynamic Set k based on sequence length

t Minimum amount of tokens in order to drop tokens

- Base models: BERT, RoBERTa
 - 1 epoch
 - $k \{1, 2, 3\}$
 - $t \{8, 10, 12\}$
 - *dynamic*
- } *static*
- Learning rate: $\{3e-5, 1e-5\}$
 - Batch size: $\{64, 128, 256, 512\}$

PLM	AA	STS				Transfer	
		k	t	BS	LR	BS	LR
BERT	naive	1	10	64	3e-5	64	1e-5
	rollout	dynamic	8	64	3e-5	512	1e-5
RoBERTa	naive	dynamic	12	64	1e-5	256	1e-5
	rollout	dynamic	12	64	1e-5	128	1e-5

Model	STS12	STS13	STS14	STS15	STS16	STS-B	SICK-R	Avg.
SimCSE-BERT _{base} ♥	68.40	82.41	74.38	80.91	78.56	76.85	72.23	76.25
+ ADD _{naive}	71.00	82.24	75.10	82.73	79.03	78.51	72.12	77.25
+ ADD _{rollout}	65.20	77.98	71.26	80.62	77.27	76.26	69.68	74.04
SimCSE-RoBERTa _{base} ♥	70.16	81.77	73.24	81.36	80.65	80.22	68.56	76.57
+ ADD _{naive}	67.45	83.43	74.67	82.48	81.69	82.00	70.43	77.45
+ ADD _{rollout}	65.34	80.97	71.29	81.08	80.34	79.83	69.54	75.48

STS task performance for sentence embeddings (Spearman's correlation, "all" setting).
 The best performance for the corresponding task is marked in bold, the second best is in italics.
 ♥: results from Gao et al. (2021); other results are evaluated by us.

Model	MR	CR	SUBJ	MPQA	SST	TREC	MRPC	Avg.
SimCSE-BERT _{base} ♥	81.18	86.46	94.45	88.88	85.50	89.80	74.43	85.81
+ MLM ♥	82.92	87.23	95.71	88.73	86.81	87.01	78.07	86.64
+ ADD _{naive}	81.82	86.89	94.83	89.43	85.28	89.40	75.25	86.13
+ MLM	82.18	87.74	95.66	88.16	86.55	91.00	75.07	86.62
+ ADD _{rollout}	81.41	85.72	94.79	89.32	84.84	88.60	75.07	85.68
+ MLM	82.40	87.97	95.62	89.38	86.93	91.20	75.59	87.01
SimCSE-RoBERTa _{base} ♥	81.04	87.74	93.28	86.94	86.60	84.60	73.68	84.84
+ MLM ♥	83.37	87.76	95.05	87.16	89.02	90.80	75.13	86.90
+ ADD _{naive}	82.30	88.05	93.70	87.50	88.25	84.60	74.84	85.61
+ MLM	83.86	89.06	94.65	88.27	89.51	90.60	76.75	87.53
+ ADD _{rollout}	82.08	88.40	93.13	87.54	87.97	87.00	75.88	86.00
+ MLM	84.68	89.91	94.97	88.37	90.61	92.20	78.43	88.45

Transfer task performance for sentence embeddings, measures represent accuracy.

The best performance for the corresponding task is marked in bold, the second best is in italics.

♥: results from Gao et al. (2021); other results are evaluated by us.

MLM: MLM is added as an auxiliary task with $\lambda = 0.1$.

- Quantification of token-relevance by **simple Attention aggregation**
- Overall improvement of performance
- Applicable to **any self-contrastive network**

Future Directions

- Explainability in Language Models
 - Identify important words/tokens in a sentence
 - Reducing sentences to essential tokens

Thank you for your $\text{softmax} \left(\frac{QK^T}{\sqrt{d_k}} \right) V!$

Find the paper on



github.com/fstermann/attention-driven-dropout

Find me on



github.com/fstermann



linkedin.com/in/fstermann



fabian.stermann@convalid-ai.de